

Artificial Locomotion

Should we **LEARN** or **OPTIMIZE**?

Nicolas Mansard

Gepetto
LAAS-CNRS & ANITI



Legged Locomotion



Robots that can **go anywhere**.

Inspirations ...



Marc Raibert et coll. MIT Leg Lab (... - 1995)



Karl Sims et coll. Evolved Virtual Creatures (1994)



Why is legged robotics *particularly* interesting?



Task-based morphology

wheels



On flat floor

4 legs



If you have enough space

6 legs



Increased safety

2 legs



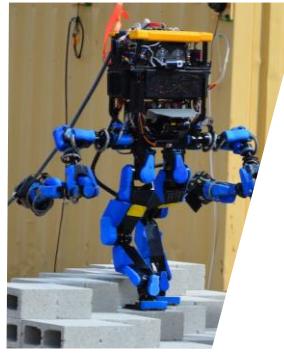
Constrained range

X legs+arms+wheels



Mixed solution

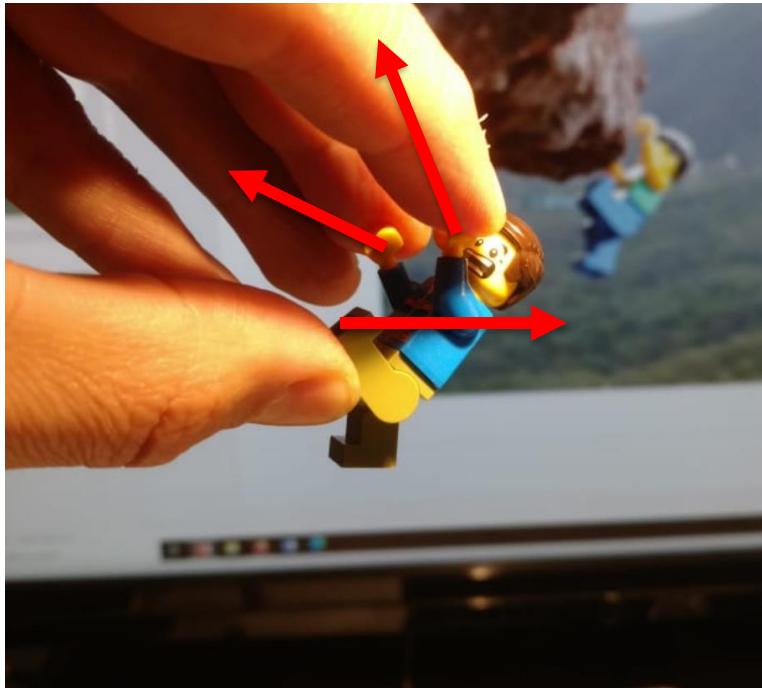
???



Move in human-centric environments ...



Locomotion and *manipulation* are dual problems



Manipulation



Locomotion

Apply forces on your environment ...



Versatile manufacturing



Flying hands

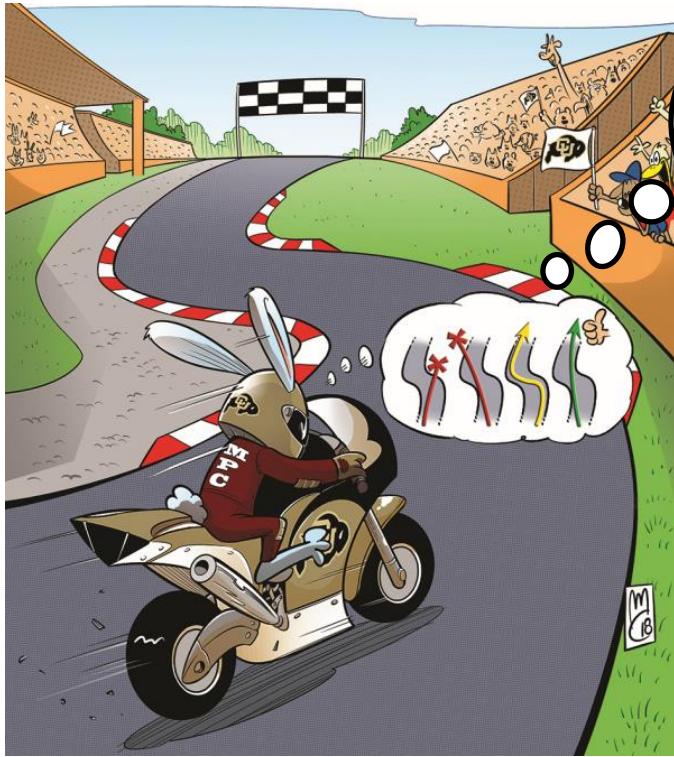


Mobile manipulators



Predictive control

Original artwork by Michele Carminati,
commissioned by Marco M. Nicotra (U. Colorado Boulder)



Decide: future robot trajectory

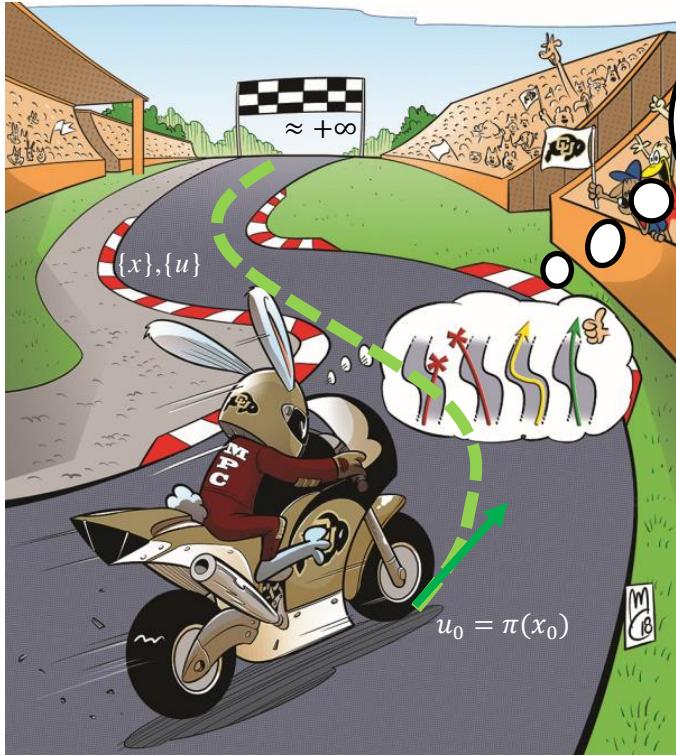
By optimizing an objective function
(eg *minimum energy*)

Imposing:

- Known initial state
- Known evolution model (simulator)
- ... and other constraints

Predictive control

Original artwork by Michele Carminati,
commissioned by Marco M. Nicotra (U. Colorado Boulder)



$$\min_{\substack{X=(Q,\dot{Q}), \\ U=\tau}} \int_0^T l(x_t, u_t) dt$$

so that $x(0) = \hat{x}$
 $\dot{x}(t) = f(x(t), u(t)), \forall t=0..T$

Simulator

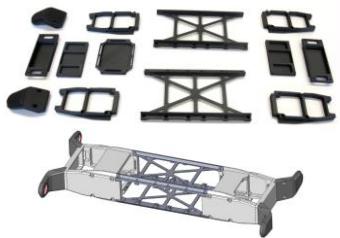
Example



Core Components



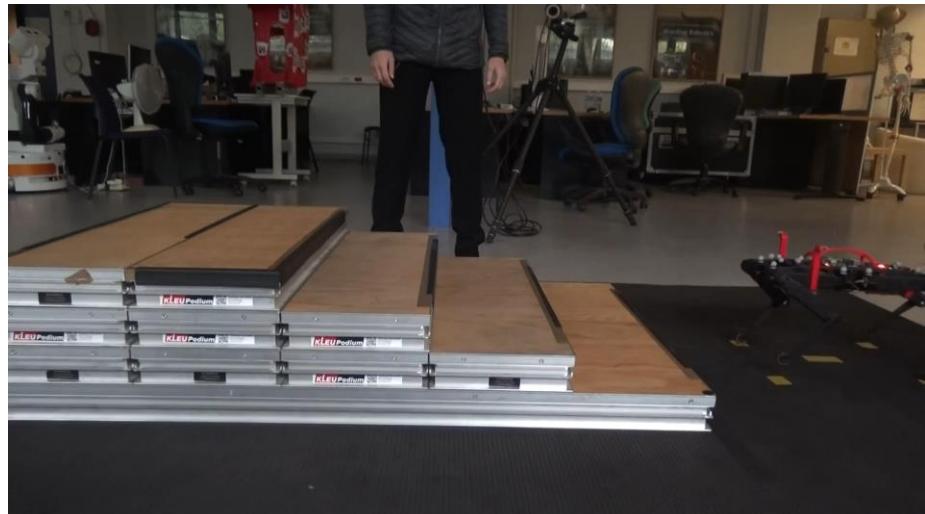
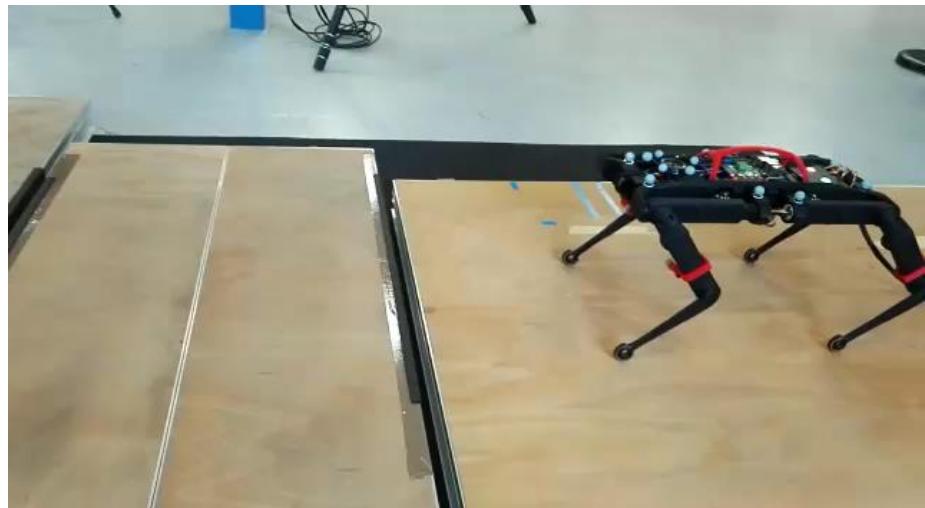
3D Printed Actuator Shells



3D Printed Body Structure

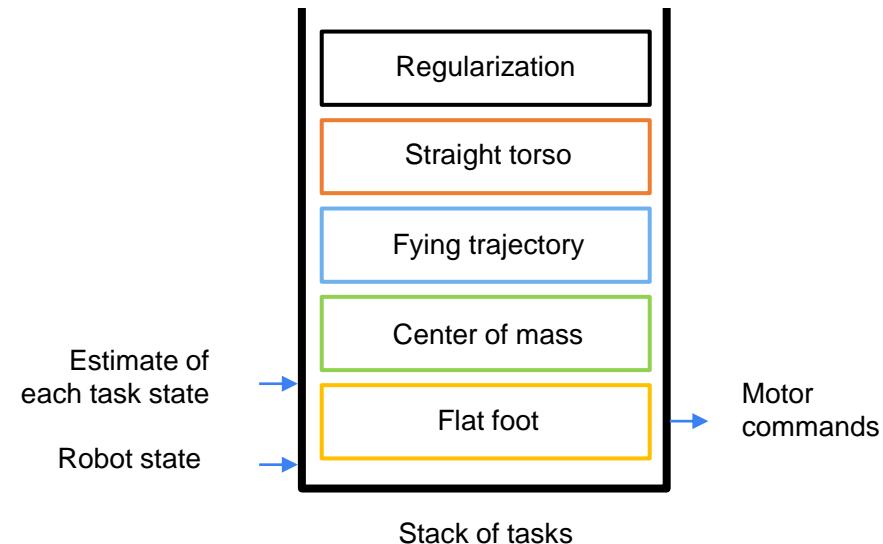
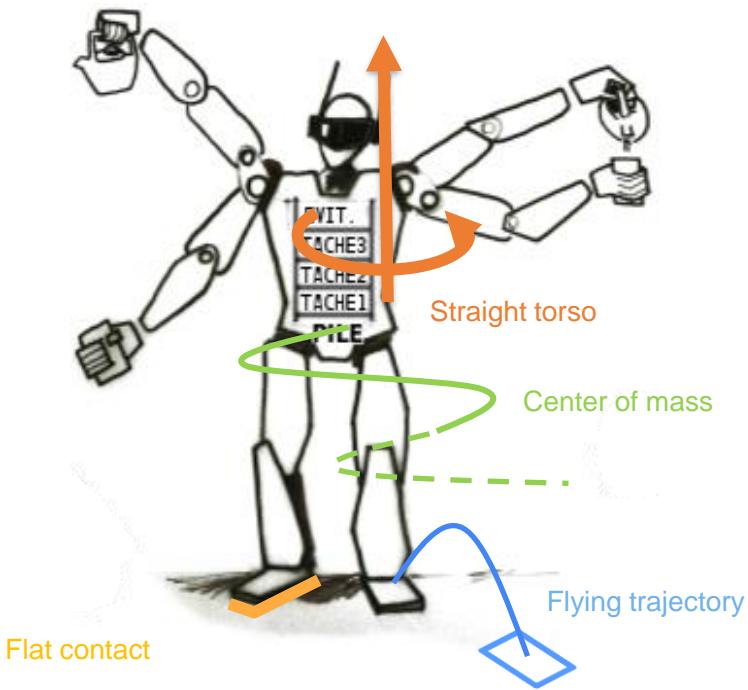


Custom Electronics

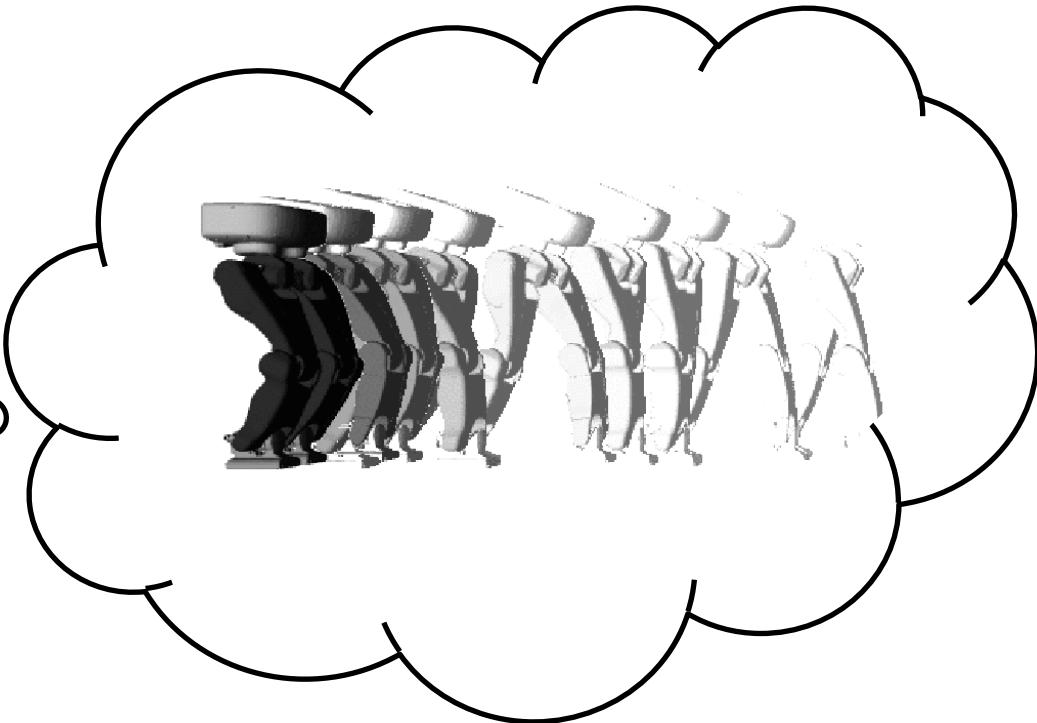
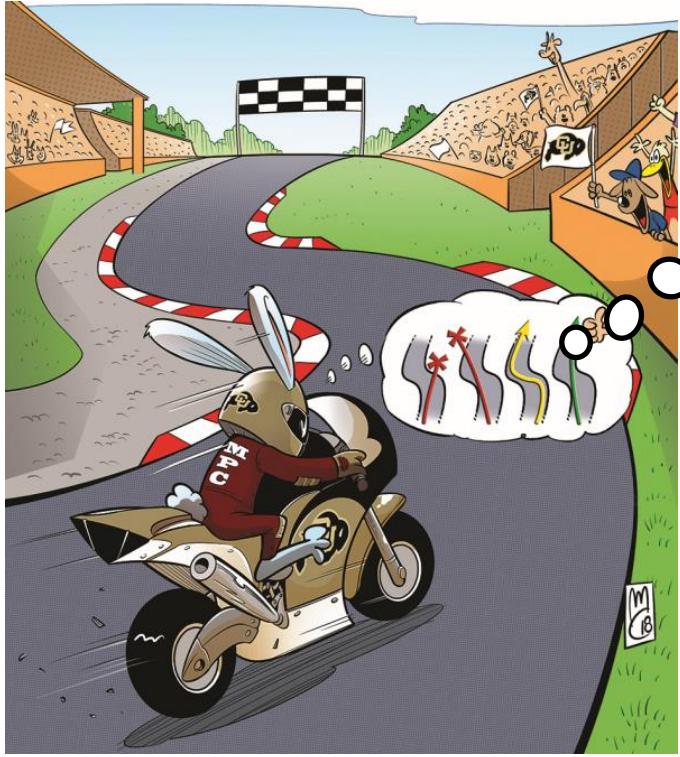


Operation control

... or *task-space control*



Original artwork by Michele Carminati,
commissioned by Marco M. Nicotra (U. Colorado Boulder)

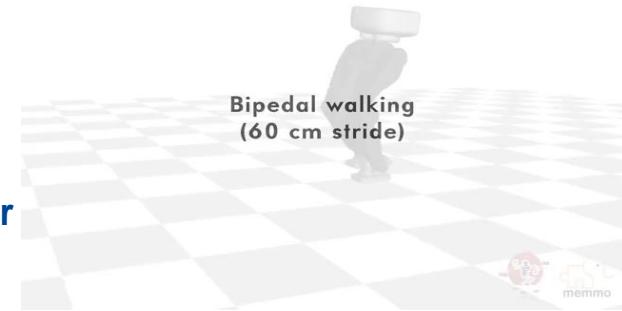


Efficient solvers ...



Carlos Mastalli
Univ. Watt @ Edinburgh

- Features expected from a good optimal control solver
 - Stable prediction: **multiple shooting**
 - Sparsity: **differential dynamic programming**
 - Strict constraints: **augmented Lagrangian**
 - Our solver incorporates all three !
- Performance on real case studies
 - 4 trot cycles for a quadruped: 8K vars, 12 iterations, **9ms / iter**
 - 2 steps for a humanoid: 12K vars, 18 iterations, **13ms / iter**



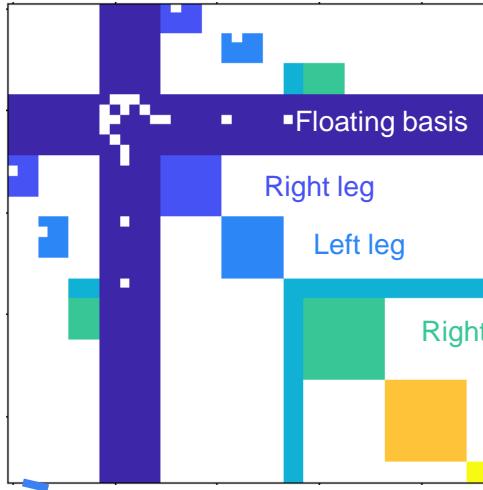
Crocoddyl
<https://github.com/loci-3d/crocoddyl>

BSD
BSD-2 License

... for efficient problems



Justin Carpentier
Inria Paris PR[AI]RIE



BSD
BSD-2 License



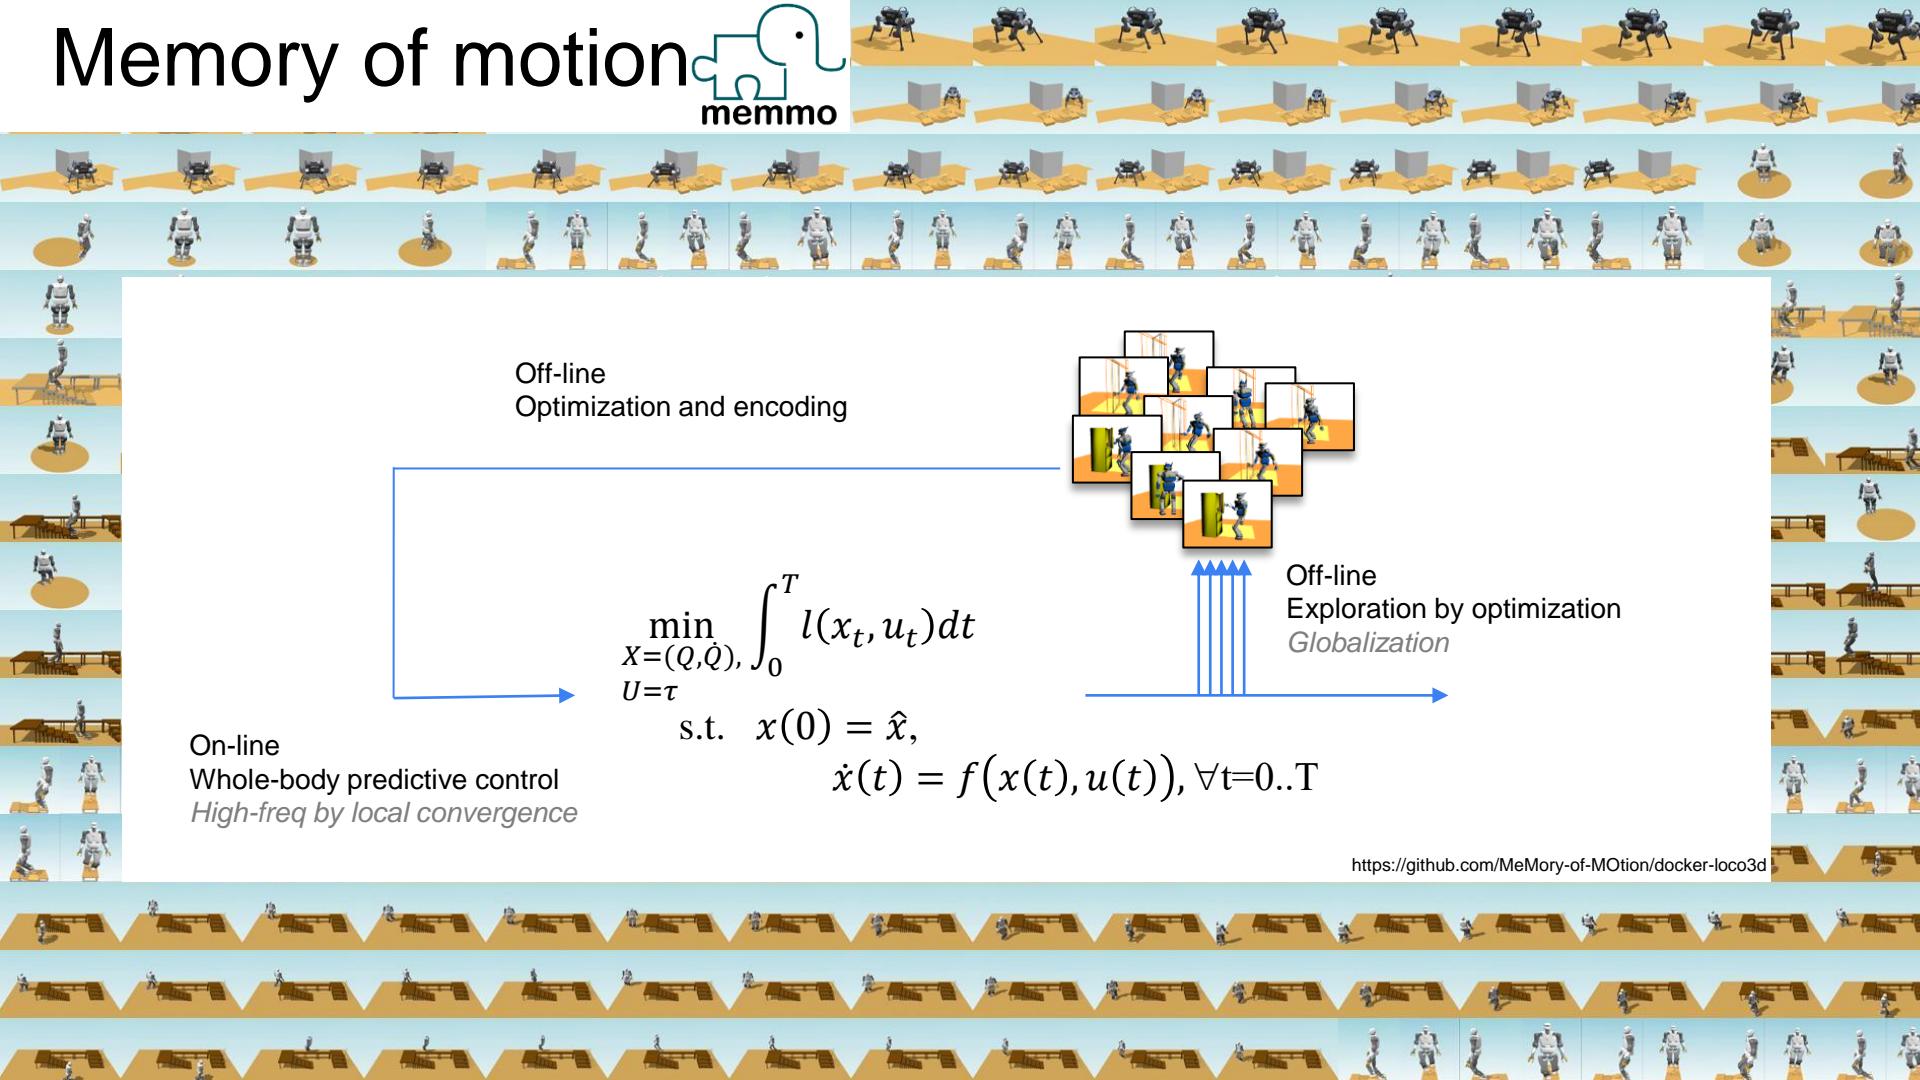
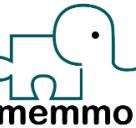
Pinocchio

<https://github.com/stack-of-tasks/pinocchio>

$$\begin{bmatrix} L_{xx} & & & & L_{xu} & & -I & F_x^T \\ & L_{xx} & & & L_{xu} & & -I & F_x^T \\ & & L_{xx} & & L_{xu} & & -I & -I \\ L_{ux} & & & L_{uu} & & F_u^T & & F_u^T \\ & L_{ux} & & & L_{uu} & & & \\ & & -I & & & & & \\ & & F_x & & & & & \end{bmatrix}$$

Carpentier et al. (2019) – <https://hal.science/hal-01866228v2>

Memory of motion

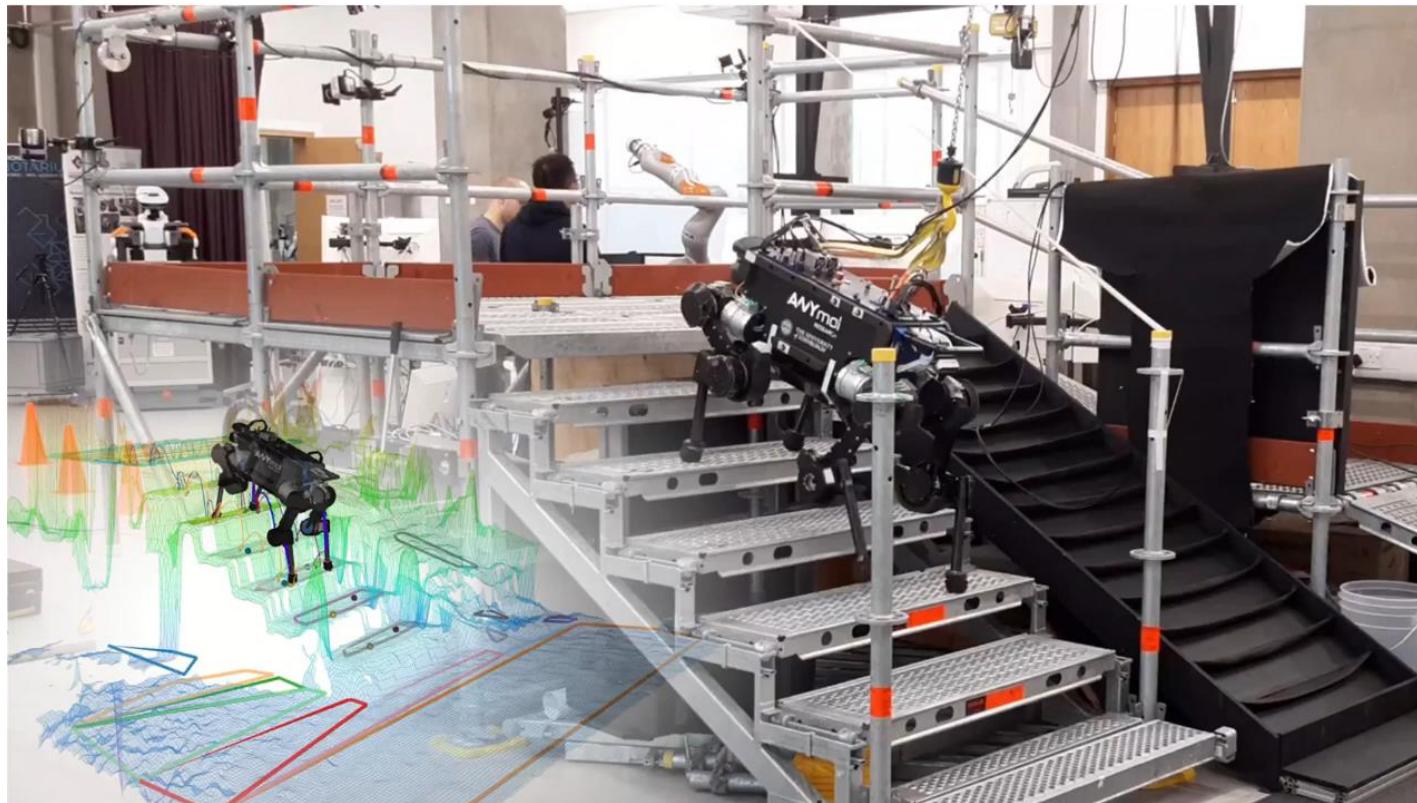


Efficient solver ... for efficient problems



Ewen Dantec

Efficient solver ... for efficient problems



Thomas Steve
Corberes Tonneau



Constraints with proximal augmented lagrangian

$$\min_x f(x) \text{ s.t. } g(x) \leq 0$$

- (Proximal) augmented Lagrangian

$$L_\mu(x, \lambda) = f(x) + \frac{1}{2\mu} \|g(x) + \lambda\|_+^2 - \frac{\mu}{2} \|\lambda\|^2$$

- Primal update

$$\begin{bmatrix} \nabla^2 L & \nabla g \\ \nabla g^T & -\mu \end{bmatrix} \begin{bmatrix} \Delta x \\ \lambda \end{bmatrix} = - \begin{bmatrix} \nabla f + \lambda^T \nabla g \\ 0 \end{bmatrix}$$

- Dual update

$$\lambda += g/\mu$$

Optimal control induces sparse structure

$$\left[\begin{array}{ccc|cc|cc} L_{xx} & & L_{xu} & -I & F_x^T \\ & \ddots & \ddots & \ddots & \ddots & \\ & L_{xx} & L_{xu} & & -I & F_x^T \\ & & L_{xx} & & & -I \\ \hline L_{ux} & & L_{uu} & F_u^T & & \\ & \ddots & \ddots & \ddots & & \\ & L_{ux} & L_{uu} & & F_u^T & \\ \hline -I & F_x & -I & & & \\ & & \ddots & & & \\ & F_x & -I & & & \\ & & & & F_u & \\ & & & & & \end{array} \right]$$

$$\left[\begin{array}{cccccc} \Gamma_1 & M_1^T & 0 & 0 & \cdots & 0 \\ M_1 & \Gamma_2 & M_2^T & 0 & \cdots & 0 \\ 0 & M_2 & \Gamma_3 & M_3^T & \cdots & 0 \\ 0 & 0 & M_3 & \Gamma_4 & \ddots & 0 \\ \vdots & \vdots & \vdots & \ddots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & \ddots & \Gamma_T \end{array} \right] \begin{bmatrix} s_1 \\ s_2 \\ s_3 \\ s_4 \\ \vdots \\ s_T \end{bmatrix} = \begin{bmatrix} g_1 \\ g_2 \\ g_3 \\ g_4 \\ \vdots \\ g_T \end{bmatrix}$$



Crocoddyl

Progress in numerical optimization

- With a augmented Lagrangian approach (with Willow@Inria)



- Using proximal formulation to handle the conditioning
- Following the main-stream literature
- Implemented in a proposition of renew of Crocoddyl (Aligator package)

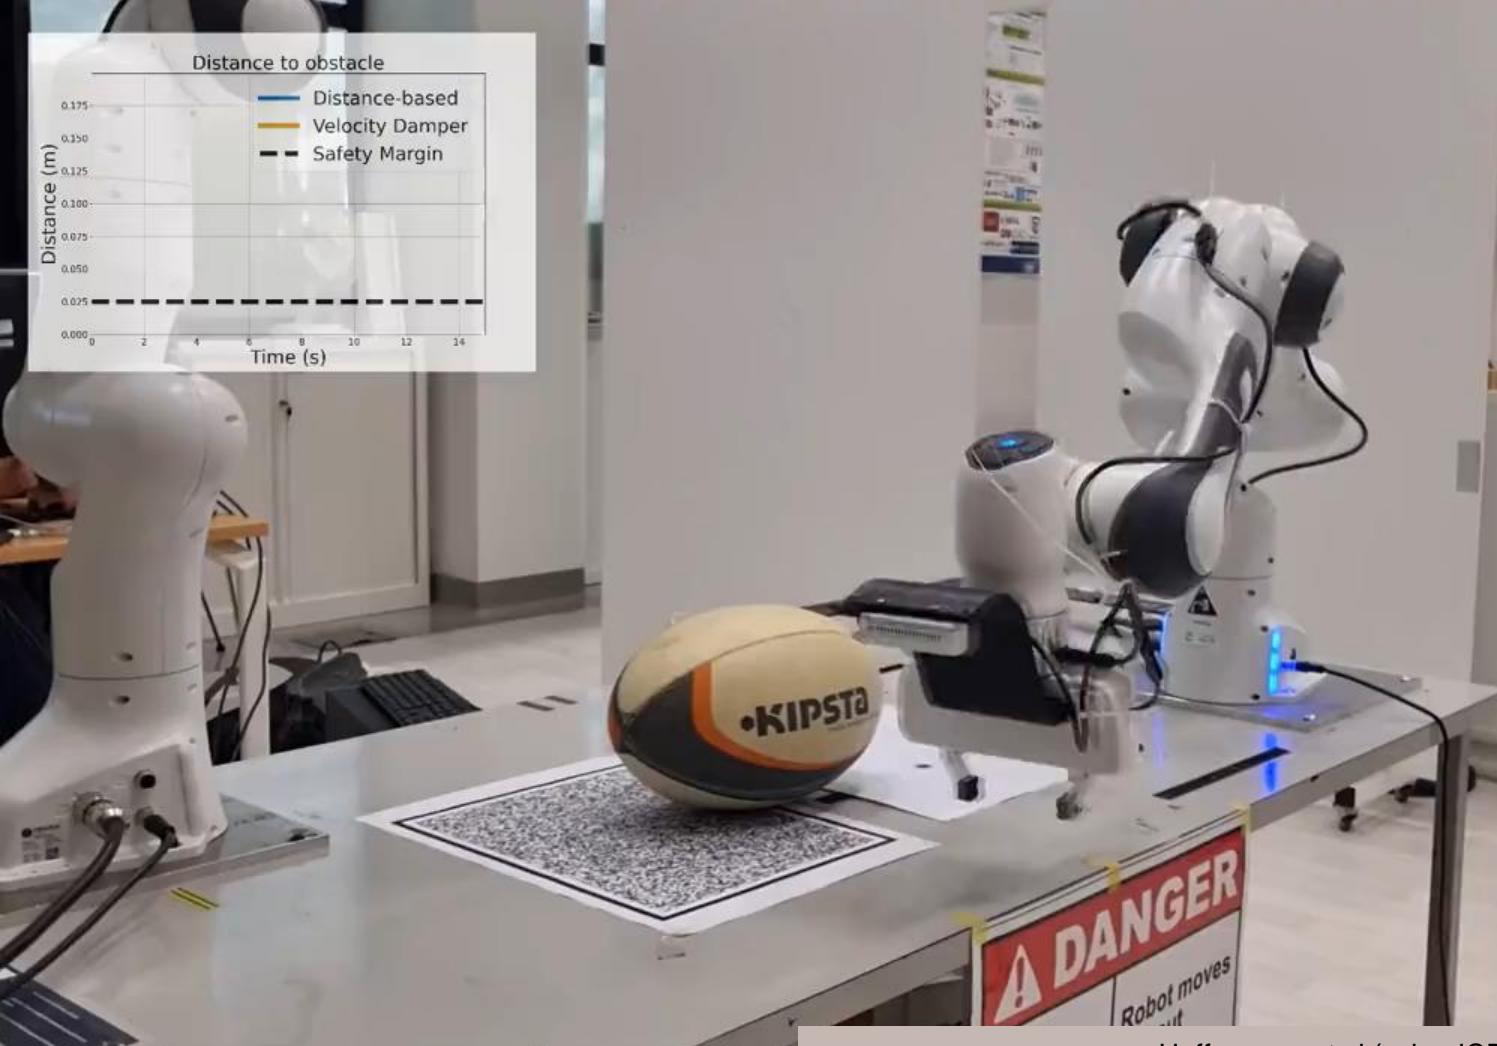
Jallet et al. (subm TRO 2023) – <https://hal.science/hal-04332348v1>

- With a SQP approach (with MiM@NYU)



- Using the operator-splitting method (OSQP) with Riccati linear solver
- Following the main-stream literature
- Implemented as an ad-on to Crocoddyl (MiM-solver package)

Jordana+Kleff+Meduri et al (subm TRO 2023) – <https://ens.hal.science/hal-04330251v1>

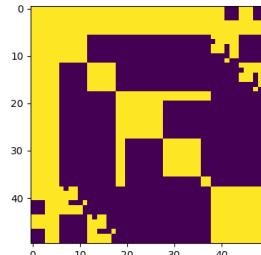


Efficient dynamics with constraints

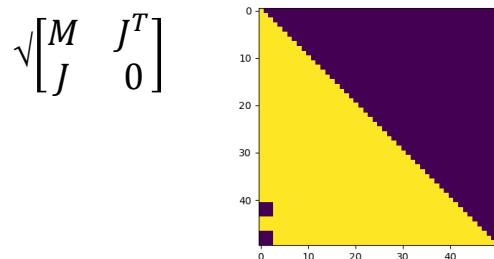
- Forward dynamics with (bilateral) constraints

$$M\ddot{q} + b = \tau + J^T f$$
$$J\ddot{q} = \gamma_0$$

- KKT structure

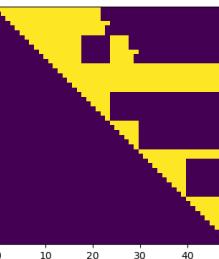


Classical Cholesky decomposition
is not sparse

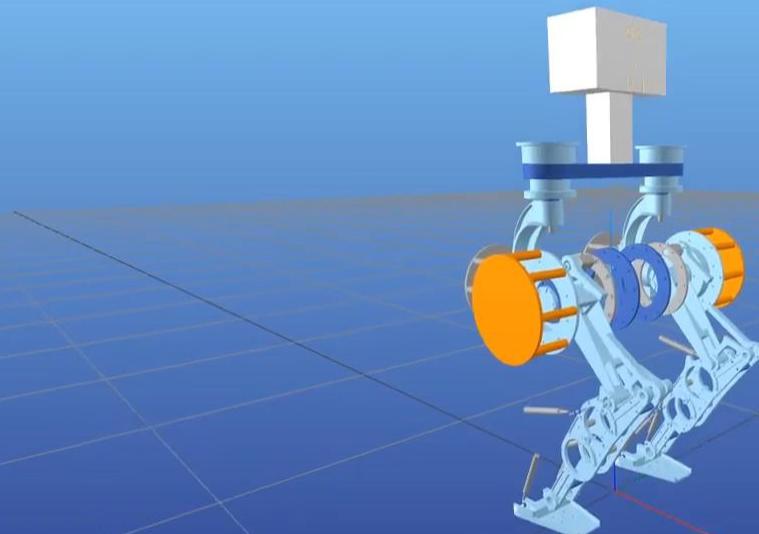


$$\sqrt{\begin{bmatrix} 0 & J \\ J^T & M \end{bmatrix}}$$

Permutting the terms and using
UDUT decomposition

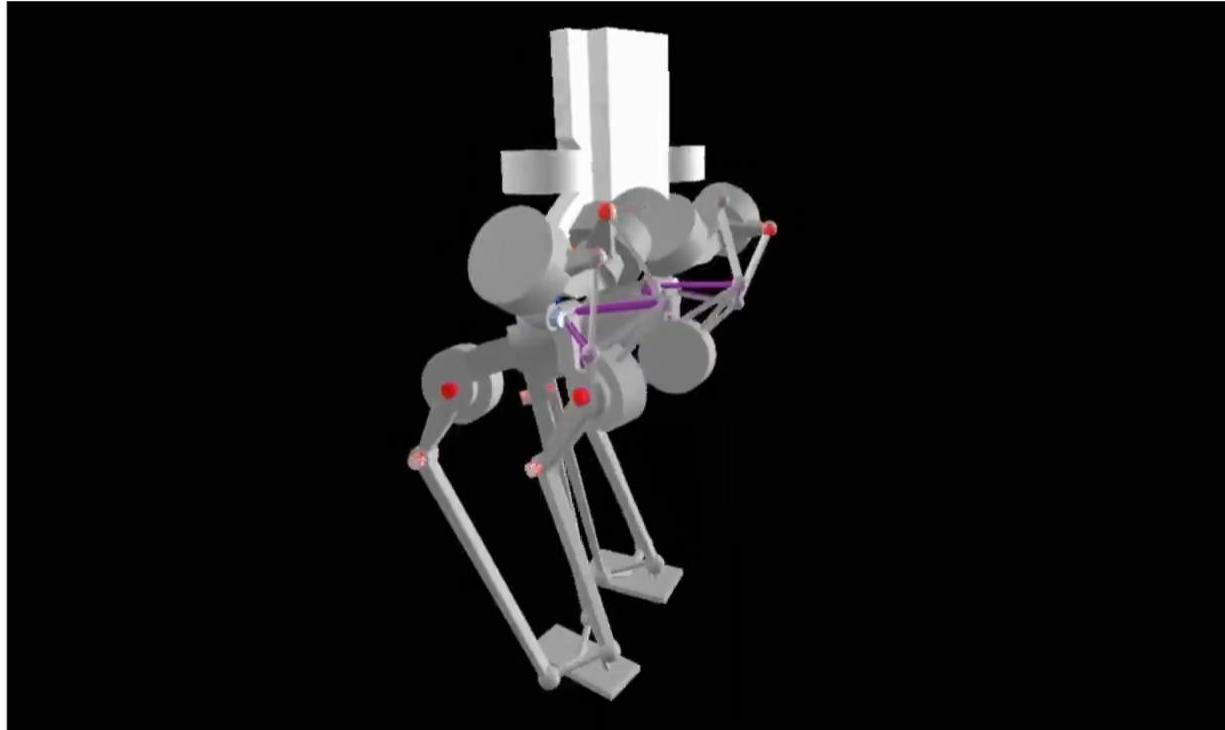


Closed Kinematic Loops as Constraints

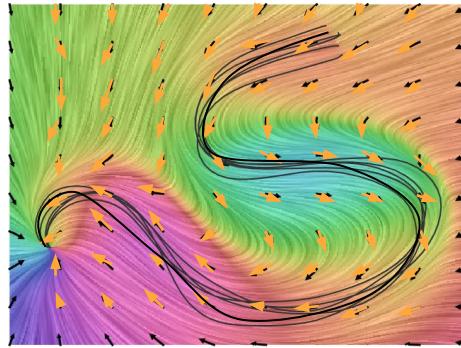


Kinematics of the architecture

New biped architecture proposed with fully-parallel kinematics



$$\begin{aligned}
& \min_{\substack{x=(Q,\dot{Q}), \\ U=\tau}} \int_0^T l(x_t, u_t) dt \\
& \text{s.t. } x(0) = \hat{x}, \\
& \quad \dot{x}(t) = f(x(t), u(t)), \forall t=0..T
\end{aligned}$$



Trajectory optimization

$$U: t \rightarrow u(t)$$

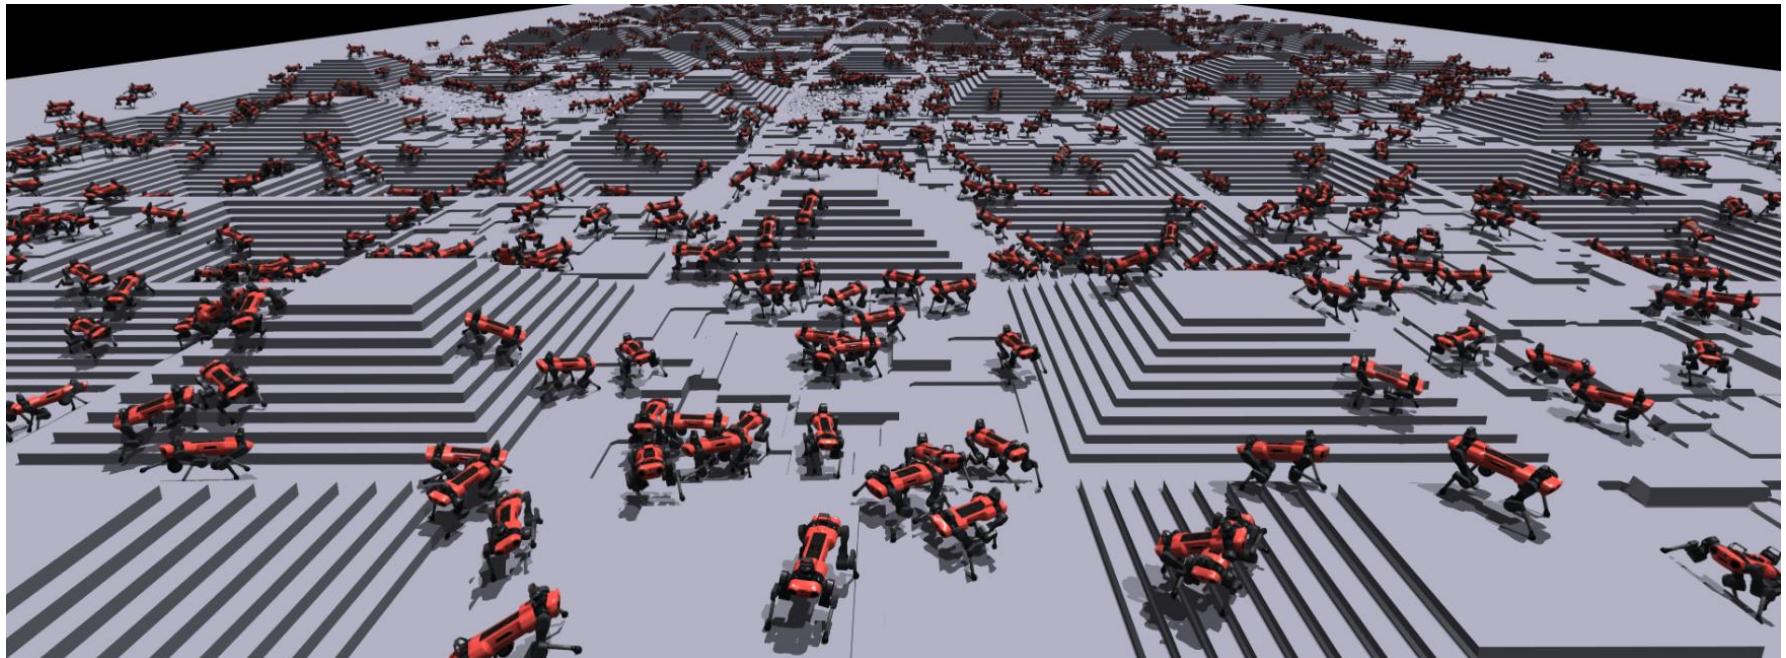
Motion planning

Policy optimization

$$\Pi: x \rightarrow u = \Pi(x)$$

Reinforcement learning

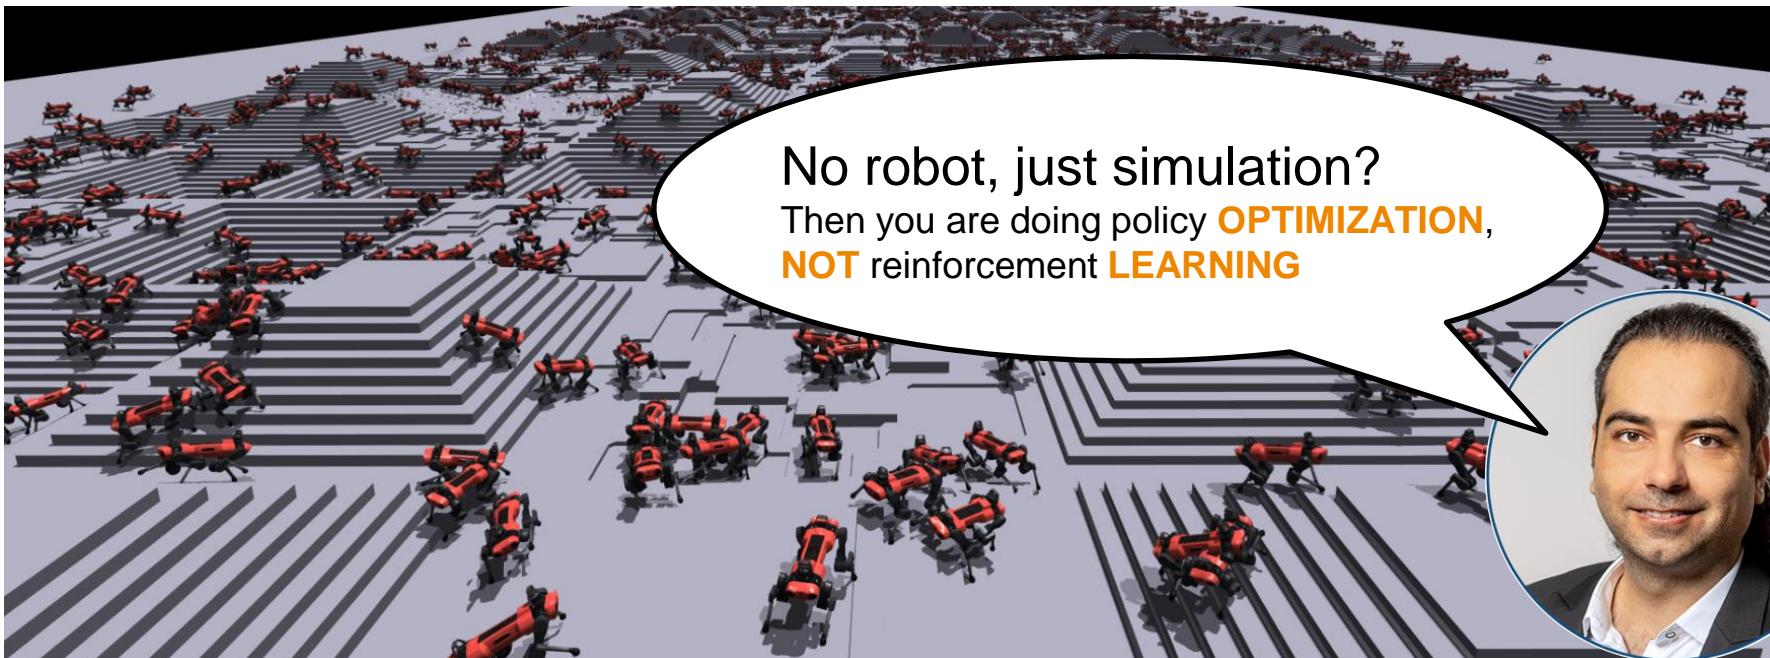
Policy optimization



Optimize the policies in physics simulators, then transfer on real robots.

Policy optimization

(aka *RL from simu*)

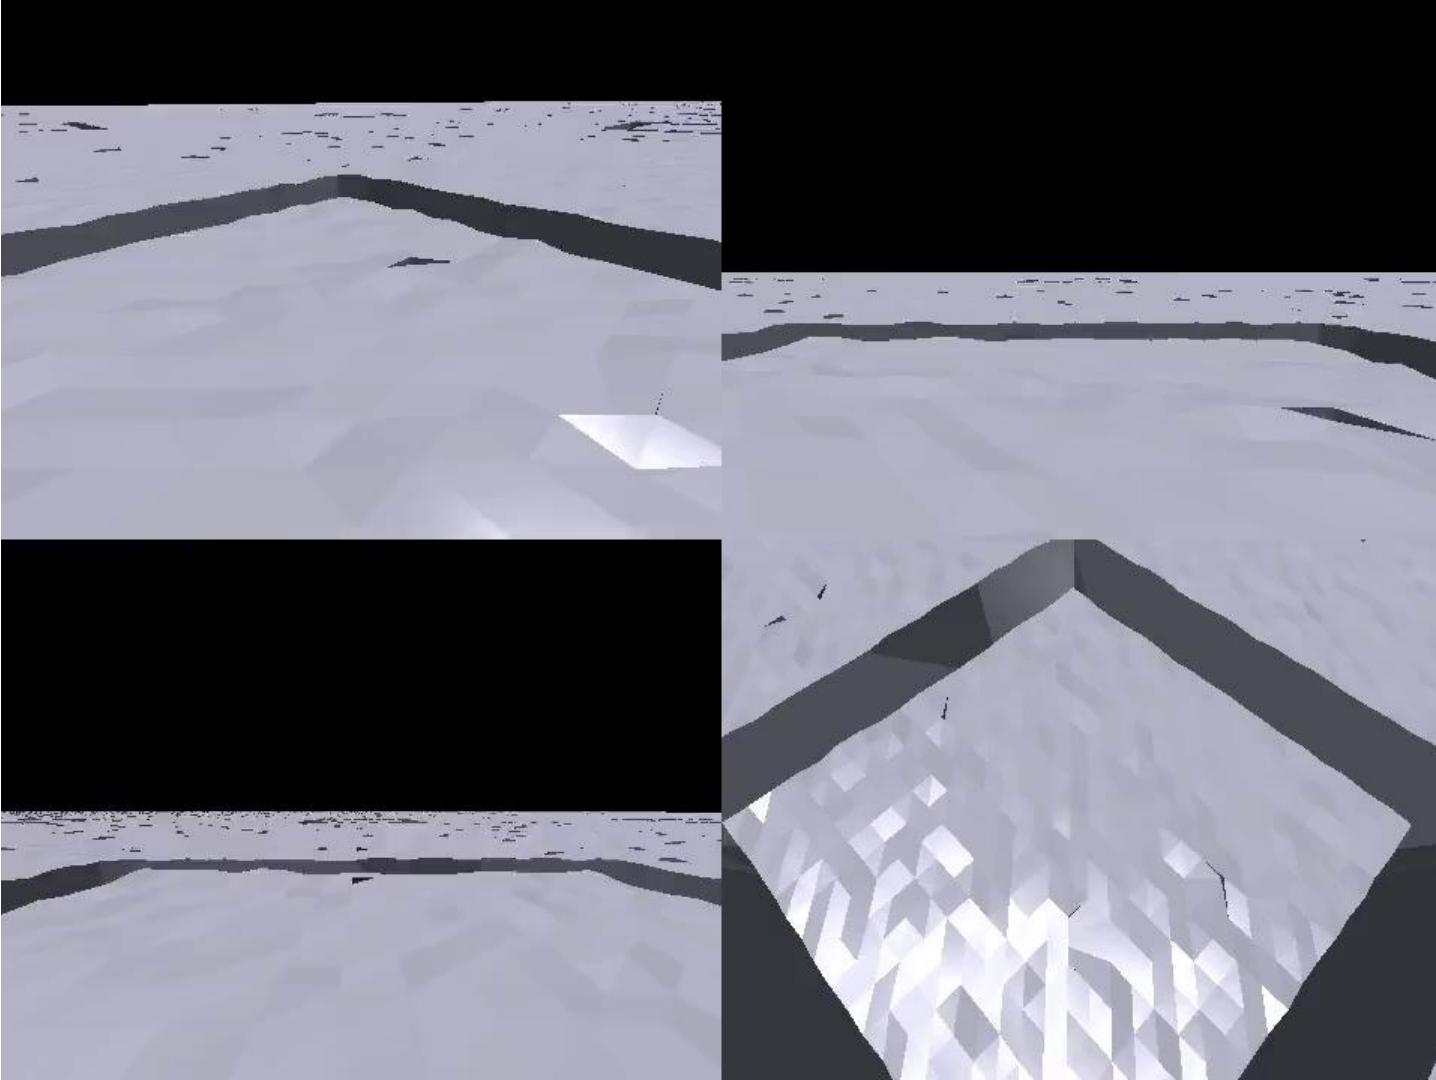


No robot, just simulation?

Then you are doing policy **OPTIMIZATION**,
NOT reinforcement **LEARNING**

Optimize the policies in physics simulators, then transfer on real robots.



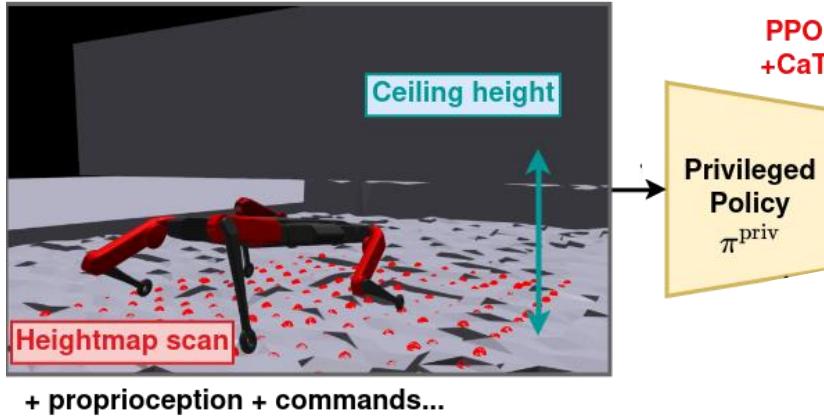


Example of shot put forbidden technique

GLIDE TECHNIQUE



Stage 1: Reinforcement Learning with Privileged Information

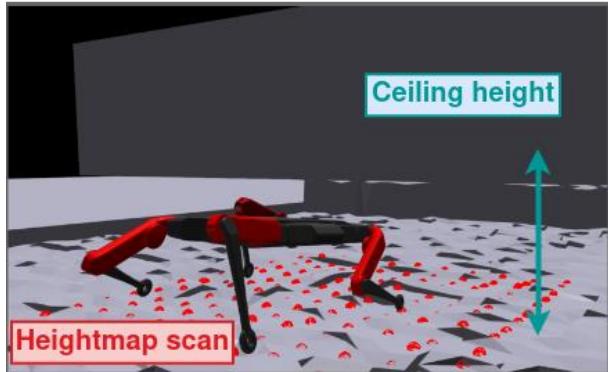


Constrained RL with CaT and PPO from **privileged information** without images:

- Elevation map
- Ceiling height

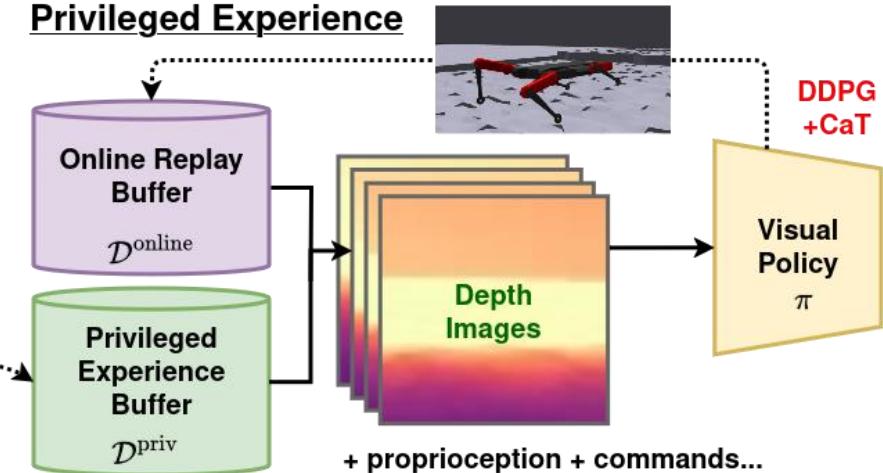
Collecting data is fast

Stage 1: Reinforcement Learning with Privileged Information



+ proprioception + commands...

Stage 2: Visuomotor Reinforcement Learning from Privileged Experience

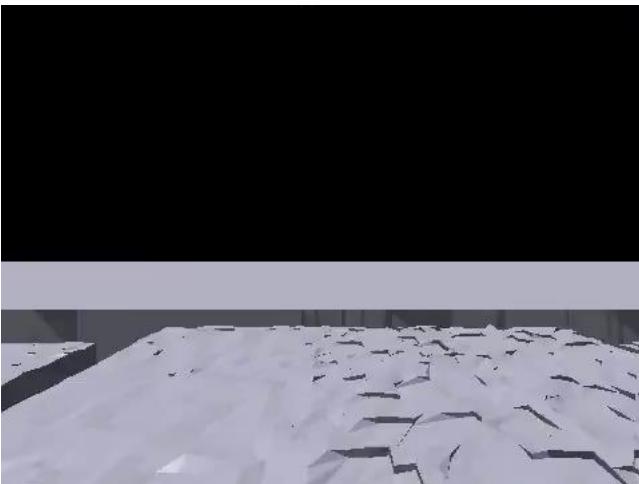
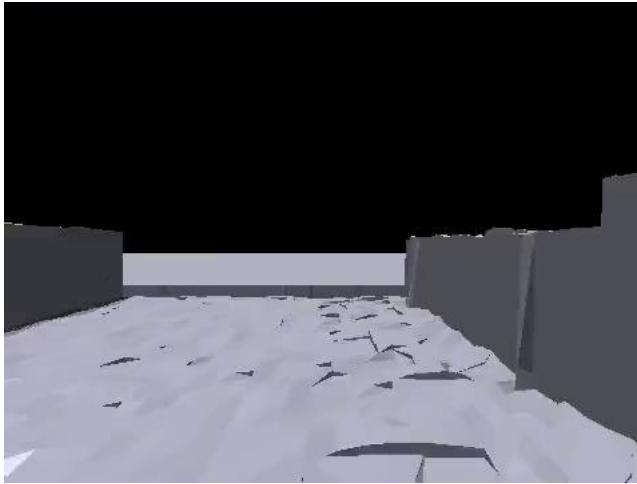
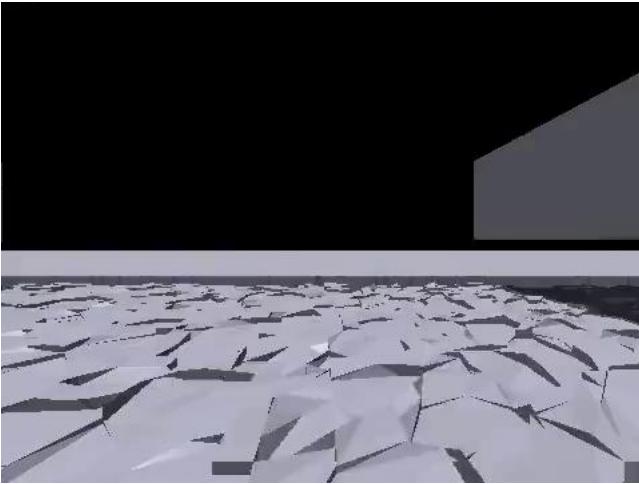


+ proprioception + commands...

Highly sample-efficient RL + CaT

- Off-policy RL
- Privileged experience transfer
- High replay ratio, asymmetric actor-critic, layer norm critic...

Collecting data is slow



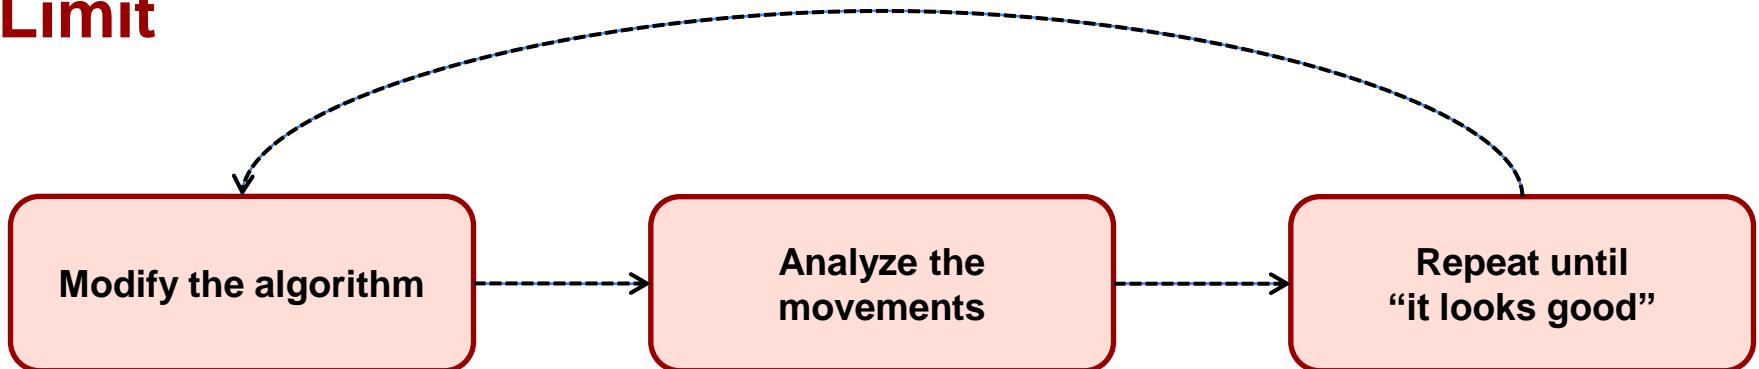
CoRL 2024

SoloParkour: Constrained Reinforcement Learning for Visual Locomotion from Privileged Experience

Elliot Chane-Sane*, Joseph Amigo*, Thomas Flayols, Ludovic Righetti, Nicolas Mansard



Limit



ex: reward engineering



Looking good = animal-like ?

Can we directly optimize for visually successful robot motions ?

**Can robots learn by watching thousands of
wild animal videos from the internet ?**

Video-Based Reward Function

Anymal Kingdom Dataset

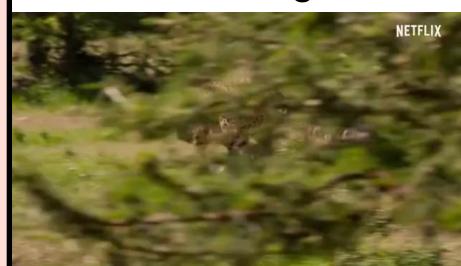
Keeping still



Walking



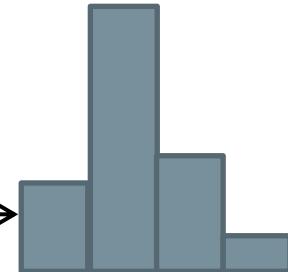
Running



Jumping



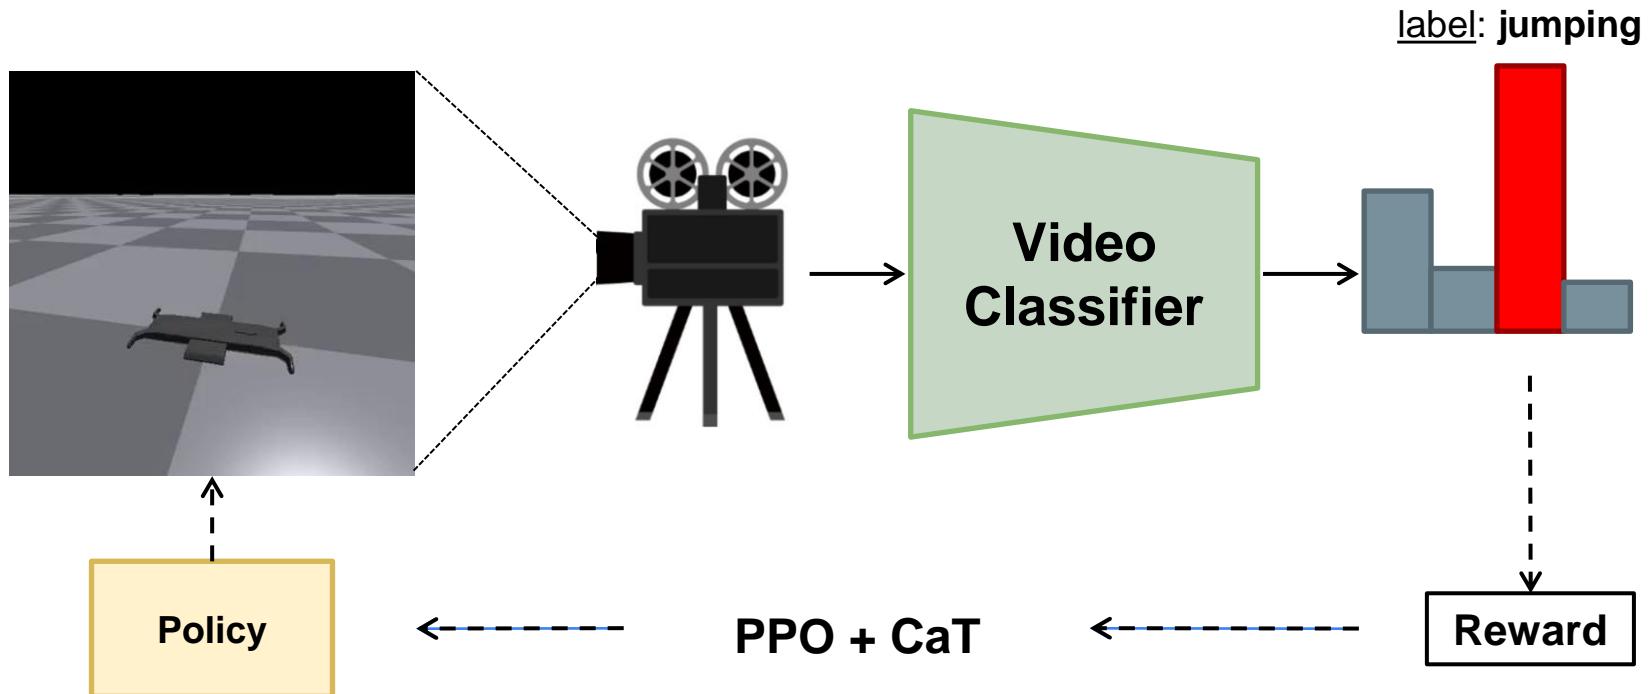
Video
Classifier
(UniFormer)



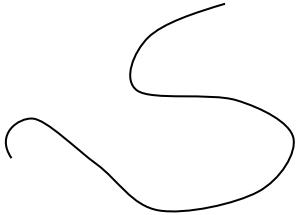
selected 8,791 videos for
training

Animal Kingdom: A Large and Diverse Dataset for Animal Behavior Understanding, Ng et al., CVPR 2022
UniFormer: Unified Transformer for Efficient Spatiotemporal Representation Learning, Li et al., ICLR 2022

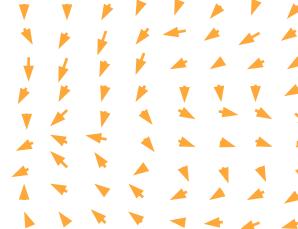
Reinforcement Learning from Wild Animal Videos



The classifier trained on animal videos transfers **zero-shot**



$$\begin{aligned} & \min_{\substack{X=(Q,\dot{Q}), \\ U=\tau}} \int_0^T l(x_t, u_t) dt \\ & \text{s.t. } x(0) = \hat{x}, \\ & \quad \dot{x}(t) = f(x(t), u(t)), \forall t=0..T \end{aligned}$$



- ✓ trajectory optimization
- ✓ super-linear convergence
- ✓ real-time computation
- ✓ constraint satisfaction
- ✗ local minima (no global policy)
- ✗ difficulty with discontinuous dynamics
- ✗ no inclusion of multi-modal sensing

- ✓ (global) policy optimization
- ✓ handles discontinuities
- ✓ multi-modal sensing inclusion
- ✗ no guaranteed convergence
- ✗ little use of model information
- ✗ difficult transfer to robots
- ✗ no constraint satisfaction

Messages to memorize



- Motion optimization is needed
 - Generalization: MPC is one-shot learning
 - Constraint satisfaction: road to IA with guarantees
 - Robotics needs a lot of numerical accuracy
- Motion learning is needed
 - Globalization comes from learning
 - No time to reinvent the wheel on-line
- There is a lot of intelligence in the hardware



Impact sociétal

Harvard computer (1950)



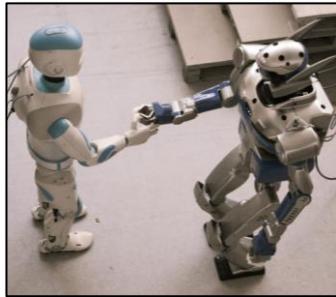
Internet



Informatique mobile



Prototypes
de laboratoire



Romeo (2010) et HRP-2 (2001)

50 ans plus tard...

?

?

Robots armés



Gadget technologique



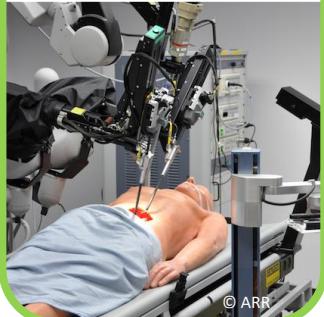
Confiscation du travail



Voiture autonome



Robotique médicale



Agriculture intelligente



Compétitivité



Libération du travail

